

Muhammad Danial Hakim Nor Azman C00253517

Changes	Time	Difficulty
Trying to create naive bayes from scratch but was stuck	1 hour	8
Cross validation for multinomial naive bayes	20 mins	2
Trying different training split to improve accuracy. Training split test , left is training and right is test accuracy , 80% = 70% , 90% = 69% , 95% = 80%	30 mins	3
Ignoring stop words like 'The' , 'And' with 80% training and get 76% accuracy	30 mins	3
Tried out using less training , achieved 82% accuracy but confusion matrix shows that the model were not able to detect positives or good mail	30 mins	3
Cross validation for Complement naive bayes	20 mins	2
Switch from MultinomialNB to ComplementNB and get 82% accuracy	45 mins	5
Change training split to 85% training and get 84% accuracy with alpha 0.3	15 mins	2
Trying new dataset on GaussianNB about wine quality where score > 6.5 is good	30 mins	3
Changing the training split with 90% training and get accuracy of 89%	20 mins	2
Confusion matrix show the model sometimes predict bad wine as good wine so try out complementNB with cross validation but get lower accuracy	40 mins	4

```
In [188... import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB, GaussianNB, ComplementNB
from sklearn.pipeline import make_pipeline
from sklearn.metrics import accuracy_score, classification_report, confusion_mat
from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt
```

```
#load data file_path = 'email_spam.csv' df = pd.read_csv(file_path) df['combined_text'] = df['title'] + ' ' + df['text'] X =
df['combined_text'] y = df['type'] #splitting the data into training and test X_train, X_test, y_train, y_test = train_test_split(X,
y, test_size=0.2, random_state=42) pipeline = make_pipeline(TfidfVectorizer(stop_words='english'),
MultinomialNB(alpha=0.3)) pipeline.fit(X_train, y_train) prediction = pipeline.predict(X_test) accuracy =
accuracy_score(y_test, prediction) print(f'Accuracy: {accuracy:.4f}') print(classification_report(y_test, prediction))
```

```
In [189... #Load data
file_path = 'email_spam.csv'
df = pd.read_csv(file_path)

df['combined_text'] = df['title'] + ' ' + df['text']

X = df['combined_text']
```

```

y = df['type']

#splitting the data into training and test
X_train ,X_test ,y_train ,y_test = train_test_split(X, y, test_size=0.15,random_

tfidf = TfidfVectorizer(stop_words='english')
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)

model = ComplementNB(alpha=0.3)

model.fit(X_train_tfidf,y_train)

predictions = model.predict(X_test_tfidf)

accuracy = accuracy_score(y_test , predictions)
print(f'Accuracy: {accuracy:.4f}')

print(classification_report(y_test,predictions))

```

Accuracy: 0.8462

	precision	recall	f1-score	support
not spam	0.80	1.00	0.89	8
spam	1.00	0.60	0.75	5
accuracy			0.85	13
macro avg	0.90	0.80	0.82	13
weighted avg	0.88	0.85	0.84	13

```

In [190... model = ComplementNB()
scores = cross_val_score(model, X_train_tfidf, y_train, cv=5, scoring="accuracy"
print(f"Cross-validation accuracy: {scores.mean():.4f} ± {scores.std():.4f}")

```


Cross-validation accuracy: 0.7619 ± 0.1436

```

In [191... cm = confusion_matrix(y_test,predictions)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negative', 'Pos
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```



```
In [192... #Load data
file_path = 'wine_data.csv'
df = pd.read_csv(file_path)
```

```
In [193... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol               1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [209... X = df.iloc[:, :-1] # exclude quality
y = (df["quality"] > 6.5).astype(int) # if score above 6.5 , its good wine and L

#splitting the data into training and test
X_train ,X_test ,y_train ,y_test = train_test_split(X, y, test_size=0.1,random_s

model = GaussianNB()

model.fit(X_train, y_train)
```

```
prediction = model.predict(X_test)

accuracy = accuracy_score(y_test , prediction)
print(f'Accuracy: {accuracy:.4f}')


print(classification_report(y_test,prediction))
```

Accuracy: 0.8938

	precision	recall	f1-score	support
0	0.97	0.91	0.94	138
1	0.58	0.82	0.68	22
accuracy			0.89	160
macro avg	0.77	0.86	0.81	160
weighted avg	0.92	0.89	0.90	160

```
In [211... cm = confusion_matrix(y_test,prediction)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negative', 'Pos
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
In [213... from sklearn import naive_bayes
import inspect

nb_classes = [cls[0] for cls in inspect.getmembers(naive_bayes, inspect.isclass)]
print(nb_classes)

['BernoulliNB', 'CategoricalNB', 'ComplementNB', 'GaussianNB', 'MultinomialNB',
 '_BaseDiscreteNB', '_BaseNB']
```

In [215...

```
model = GaussianNB()  
scores = cross_val_score(model, X_train, y_train, cv=5, scoring="accuracy")  
print(f"Cross-validation accuracy: {scores.mean():.4f} ± {scores.std():.4f}")
```

Cross-validation accuracy: 0.8443 ± 0.0167